

Data & Text Analysis

POS, NER, spacy

IMT 547 - Social Media Data Mining and Analysis

11-Feb-2021 (Week 6, Day 12)

Today's Topics

- Representing Text Data - DTM
- Lab leftover
- NER, POS
- Spacy
- Lab
- Project team
- PS2 discussion

Beautiful Soup cheatsheet

Cheatsheet

<http://akul.me/blog/2016/beautifulsoup-cheatsheet/>

Tutorial: Web Scraping with Python Using Beautiful Soup

<https://www.dataquest.io/blog/web-scraping-tutorial-python/>

Beautiful Soup documentation: <https://www.crummy.com/software/BeautifulSoup/bs4/doc/>

Representing Text Data

1. Corpus
2. Document-Term Matrix

Representing Text Data

Corpus: a collection of texts

Corpus of movie reviews: positive or negative

Unbelievably disappointing

Richly applied satire, great plot twists.

This is the greatest screwball comedy ever filmed

It was pathetic. The worst part about it was the boxing scenes

```
# Let's take a look at the updated text
data_clean = pd.DataFrame(data_df.transcript.apply(round1))
data_clean
```

		transcript
ali	ladies and gentlemen please welcome to the stage ali wong hi hello welcome thank you thank you for coming hello hello we are gonna have to get thi...	
anthony	thank you thank you thank you san francisco thank you so much so good to be here people were surprised when i told 'em i was gonna tape my special...	
bill	all right thank you thank you very much thank you thank you thank you how are you what's going on thank you it's a pleasure to be here in the gre...	
bo	bo what old macdonald had a farm e i e i o and on that farm he had a pig e i e i o here a snort there a old macdonald had a farm e i e i o this i...	
dave	this is dave he tells dirty jokes for a living that stare is where most of his hard work happens it signifies a profound train of thought the alch...	
hasan	what's up davis what's up i'm home i had to bring it back here netflix said "where do you want to do the special la chicago new york" i was like...	
jim	ladies and gentlemen please welcome to the stage mr jim jefferies hello sit down sit down sit down sit down sit down thank you boston i appre...	
joe	ladies and gentlemen welcome joe rogan what the fuck is going on san francisco thanks for coming i appreciate it god damn put your phone down ...	
john	armed with boyish charm and a sharp wit the former "snl" writer john mulaney offers sly takes on marriage his beef with babies and the time he met...	
louis	intro\nfade the music out let's roll hold there lights do the lights thank you thank you very much i appreciate that i don't necessarily agree wit...	
mike	wow hey thank you thanks thank you guys hey seattle nice to see you look at this look at us we're here this is crazy it's insane so about five yea...	
ricky	hello hello how you doing great thank you wow calm down shut the fuck up thank you what a lovely welcome i'm gonna try my hardest tonight you're t...	

From lab last class

Representing Text Data

Document-Term Matrix

<i>Tokens</i>	“<i>great</i>”	“<i>worst</i>”	“<i>plot</i>”
Unbelievably disappointing. Worst worst movie ever.		2		
Richly applied satire, great plot twists.				
This is such a great comedy				
It was pathetic. The worst part about it was the boxing scenes				
.....				

Representing Text Data

Document-Term Matrix

How to create one?

- 1. Clean text
- 2. Tokenize text
- 3. Document-Term Matrix: Put into a matrix so a machine can read it

	<i>Tokens</i>			
	<i>“great”</i>	<i>“worst”</i>	<i>“plot”</i>	<i>....</i>
Unbelievably disappointing. Worst worst movie ever.		2		
Richly applied satire, great plot twists.				
This is such a great comedy				
It was pathetic. The worst part about it was the boxing scenes				

Representing Text Data

Document-Term Matrix

Step 1: data cleaning

All right, Petunia. Wish me luck out there. You will die on August 7th. 2037.

How would a computer read this? It needs to be **cleaned first**

Common data cleaning steps that we can do on this data (using regular expressions, **python re**)

- Remove punctuation
- Remove numbers
- Lowercase letters

all right petunia wish me luck out there you will die on august

Representing Text Data

Document-Term Matrix

Step 2: tokenize the data

all right petunia wish me luck out there you will die on august **Cleaned data**

How would a computer read this? It needs to be cleaned first, **then tokenized**

What's tokenization?

Split text into smaller pieces (or tokens). The most common token size is a word. It can also be sentence, bigrams, trigrams

all right
luck out
die

Now you can do **some more cleaning** — removing stop words (which carry very little meaning)

Representing Text Data

Document-Term Matrix

Step 2: tokenize the data, followed by some more cleaning

all right petunia wish me luck out there you will die on august

How would a computer read this? It needs to be cleaned first, **then tokenized**

right petunia wish much die august

Bag of words model: This representation of the text is called a **bag of words model**.
Simple format that ignores order

Representing Text Data

Document-Term Matrix

Step 3: put into a matrix

all right petunia wish me luck out there you will die on august

How would a computer read this? It needs to be cleaned first, then tokenized, and **put into matrix**

	aaaaah	aaaaahhhhhhh	aaaaauugghhhhhh	aaaahhhh	aaah	aah	abc	abcs	ability	abject	...	zee	zen
ali	0	0	0	0	0	0	1	0	0	0	...	0	0
anthony	0	0	0	0	0	0	0	0	0	0	...	0	0
bill	1	0	0	0	0	0	0	1	0	0	...	0	0
bo	0	1	1	1	0	0	0	0	1	0	...	0	0
dave	0	0	0	0	1	0	0	0	0	0	...	0	0
hasan	0	0	0	0	0	0	0	0	0	0	...	2	1
jim	0	0	0	0	0	0	0	0	0	0	...	0	0
joe	0	0	0	0	0	0	0	0	0	0	...	0	0
john	0	0	0	0	0	0	0	0	0	0	...	0	0
louis	0	0	0	0	0	3	0	0	0	0	...	0	0
mike	0	0	0	0	0	0	0	0	0	0	...	0	0
ricky	0	0	0	0	0	0	0	0	1	1	...	0	0

Representing Text Data

Document term matrix:

- Each row is a different document
- Each column is a different term (usually words, can be bigrams etc.)
- The values are word counts

	aaaaah	aaaaahhhhhh	aaaaauugghhhhhh	aaaahhhh	aaah	aah	abc	abcs	ability	ak
ali	0	0	0	0	0	0	1	0	0	
anthony	0	0	0	0	0	0	0	0	0	
bill	1	0	0	0	0	0	0	1	0	
bo	0	1	1	1	0	0	0	0	1	
dave	0	0	0	0	1	0	0	0	0	
hasan	0	0	0	0	0	0	0	0	0	
jim	0	0	0	0	0	0	0	0	0	
joe	0	0	0	0	0	0	0	0	0	
john	0	0	0	0	0	0	0	0	0	
louis	0	0	0	0	0	3	0	0	0	
mike	0	0	0	0	0	0	0	0	0	
ricky	0	0	0	0	0	0	0	0	1	

How do you do it programmatically?

```
# We are going to create a document-term matrix using CountVec
from sklearn.feature_extraction.text import CountVectorizer

cv = CountVectorizer(stop_words='english') #first instantiate
data_cv = cv.fit_transform(data_clean_documents)
```

Leftover lab

Getting simple structured information from text

Named Entity Recognition

POS

Goal of Information Extraction (IE) systems:

- Find and understand relevant parts of text
- Gather information from many pieces of text

GOAL: Produce a structured representation of the relevant information

Named Entity Recognition (NER)

A named entity is a “real-world object” that’s assigned a name – for example, a person, a country,....

Apple **ORG** is looking at buying U.K. **GPE** startup for \$1 billion **MONEY**

NER - find and classify entities in a text

TEXT	START	END	LABEL	DESCRIPTION
Apple	0	5	ORG	Companies, agencies, institutions.
U.K.	27	31	GPE	Geopolitical entity, i.e. countries, cities, states.
\$1 billion	44	54	MONEY	Monetary values, including unit.

Hovy Text Analysis in python for social scientists, Chapter 2.4

Named Entity Recognition (NER)

Lots of real-world application


× 🔍

[All](#) [News](#) [Images](#) [Maps](#) [Videos](#) [More](#) [Settings](#) [Tools](#)

About 943,000,000 results (0.89 seconds)

University of Washington / President

Ana Mari Cauce




Ana Mari Cauce is the 33rd president of the UW. A professor of psychology and member of the UW faculty since 1986, she was named president in October 2015.

× 🔍

[All](#) [News](#) [Images](#) [Videos](#) [Maps](#) [More](#) [Settings](#) [Tools](#)

About 863,000,000 results (0.69 seconds)

Ana Mari Cauce is the 33rd **president** of the **UW**. A professor of psychology and member of the **UW** faculty since 1986, she was named **president** in October 2015.



Named Entity Recognition (NER)

Can you identify the named entities in this text?

New York City on Tuesday declared a public health emergency. At least 285 people have contracted measles in the city since September, mostly in Brooklyn's Williamsburg neighborhood. The order covers four Zip codes there, Mayor Bill de Blasio (D) said Tuesday.

New York City **GPE** on Tuesday **DATE** declared a public health emergency. At least 285 **CARDINAL** people have contracted measles in the city since September **DATE**, mostly in Brooklyn **GPE**'s Williamsburg neighborhood. The order covers four **CARDINAL** Zip **PERSON** codes there, Mayor Bill de Blasio **PERSON** (D) said Tuesday **DATE**.

Parts of Speech

e.g., nouns, verbs, and adjectives (most languages have more categories than just these three)

- ADJ: adjectives. They modify nouns to specify their properties. Examples: *awesome, red, boring*
- ADV: adverbs. They modify verbs, but also serve as question markers. Examples: *quietly, where, never*
- INTJ: interjections. Exclamations of some sort. Examples: *ouch, shhh, oi*
- NOUN: nouns. Entities in the world. Examples: *book, war, shark*
- PROPN: proper nouns. Names of entities, a subclass of nouns. Examples: *Rosa, Twitter, CNN*
- VERB: full verbs. Events in the world. Examples: *codes, submitted, succeed*

Hovy Text Analysis in python for social scientists, Chapter 2.2

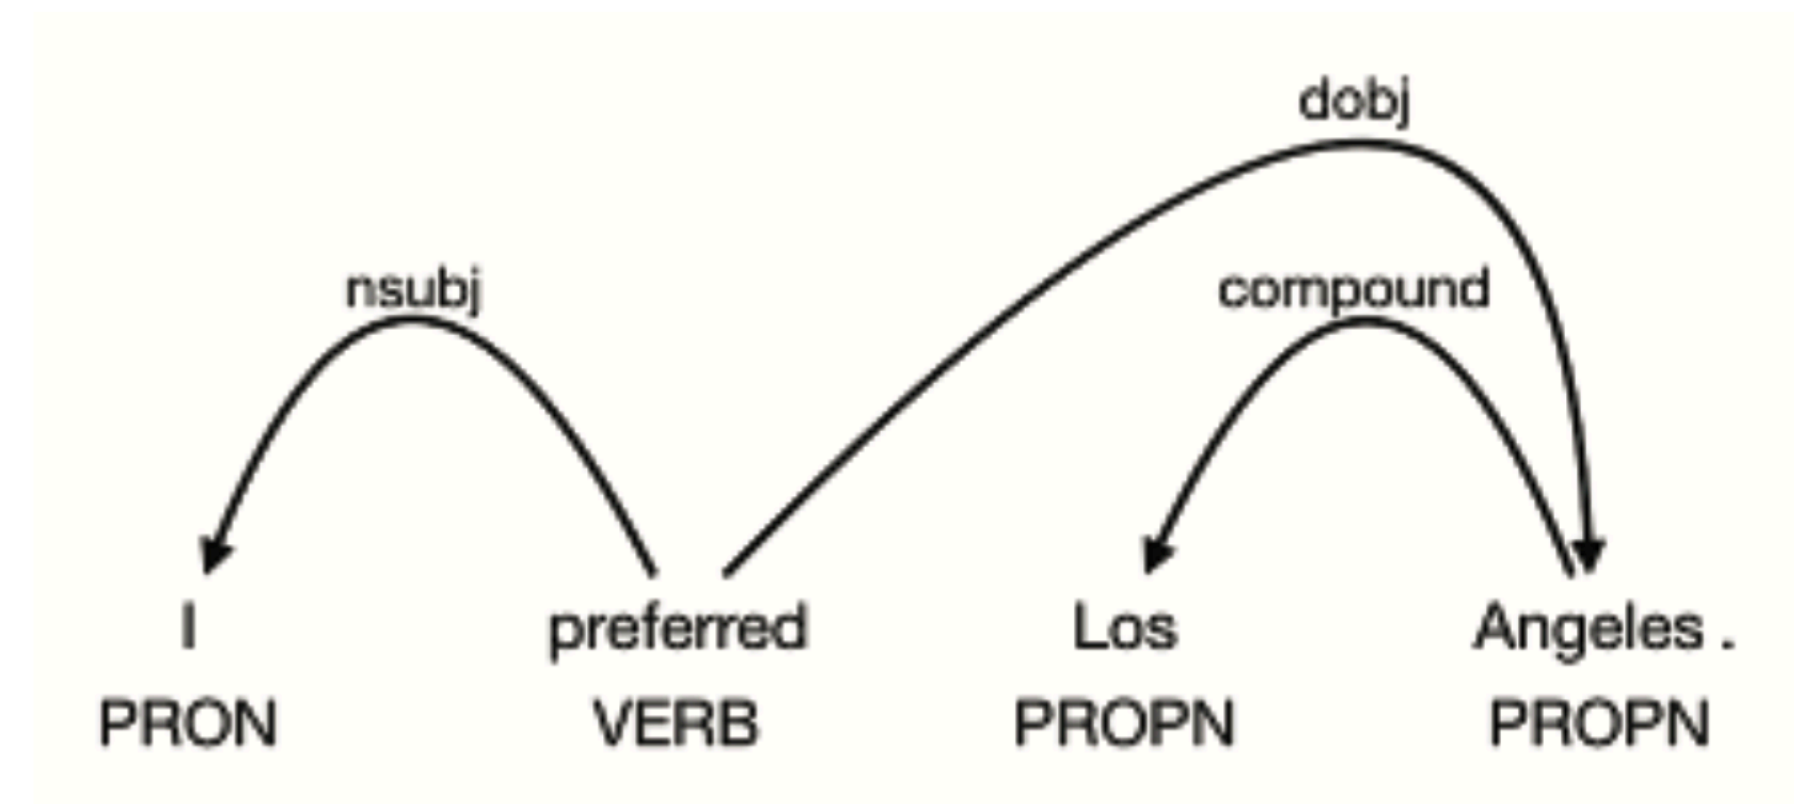


Figure 1 Example of a dependency parse.

Parts of Speech

POS in action in a DS project

Modeling Factuality Judgments in Social Media Text

Sandeep Soni

Tanushree Mitra

Eric Gilbert

Jacob Eisenstein




Figure 4: Dependency parse of an example message, with claim, source, and cue.

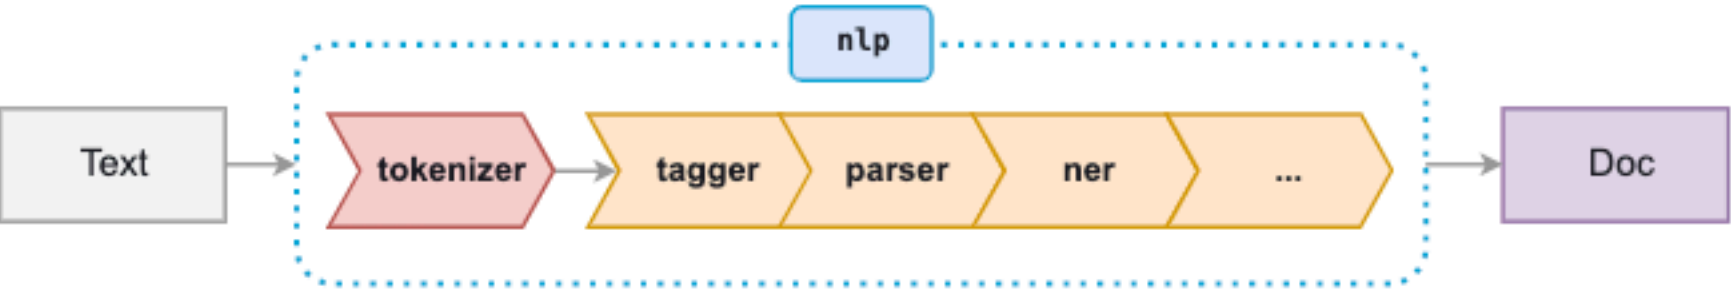
RQ: How do journalists mark quoted content as certain or uncertain, and how do readers interpret these signals?

Language processing with Spacy

<https://spacy.io/usage/processing-pipelines>

Language Processing Pipelines

When you call `nlp` on a text, spaCy first tokenizes the text to produce a `Doc` object. The `Doc` is then processed in several different steps – this is also referred to as the **processing pipeline**. The pipeline used by the [default models](#) consists of a tagger, a parser and an entity recognizer. Each pipeline component returns the processed `Doc`, which is then passed on to the next component.




NAME	COMPONENT	CREATES	DESCRIPTION
tokenizer	Tokenizer ≡	Doc	Segment text into tokens.
tagger	Tagger ≡	Doc[i].tag	Assign part-of-speech tags.
parser	DependencyParser ≡	Doc[i].head , Doc[i].dep , Doc.sents , Doc.noun_chunks	Assign dependency labels.
ner	EntityRecognizer ≡	Doc.ents , Doc[i].ent_iob , Doc[i].ent_type	Detect and label named entities.
textcat	TextCategorizer ≡	Doc.cats	Assign document labels.
...	custom components	Doc._.xxx , Token._.xxx , Span._.xxx	Assign custom attributes, methods or properties.

Part-of-speech tagging

NEEDS MODEL ⓘ

After tokenization, spaCy can **parse** and **tag** a given `Doc`. This is where the trained pipeline and its statistical models come in, which enable spaCy to **make predictions** of which tag or label most likely applies in this context. A trained component includes binary data that is produced by showing a system enough examples for it to make predictions that generalize across the language – for example, a word following “the” in English is most likely a noun.

Linguistic annotations are available as `Token attributes` . Like many NLP libraries, spaCy **encodes all strings to hash values** to reduce memory usage and improve efficiency. So to get the readable string representation of an attribute, we need to add an underscore `_` to its name:

Editable Code spaCy v3.0 · Python 3 · via Binder

```
import spacy

nlp = spacy.load("en_core_web_sm")
doc = nlp("Apple is looking at buying U.K. startup for $1 billion")

for token in doc:
    print(token.text, token.lemma_, token.pos_, token.tag_, token.dep_,
          token.shape_, token.is_alpha, token.is_stop)
```

RUN

TEXT	LEMMA	POS	TAG	DEP	SHAPE	ALPHA	STOP
Apple	apple	PROPN	NNP	nsubj	Xxxxx	True	False
is	be	AUX	VBZ	aux	xx	True	True
looking	look	VERB	VBG	ROOT	xxxx	True	False
at	at	ADP	IN	prep	xx	True	True
buying	buy	VERB	VBG	pcomp	xxxx	True	False
U.K.	u.k.	PROPN	NNP	compound	X.X.	False	False

- Text:** The original word text.
- Lemma:** The base form of the word.
- POS:** The simple UPOS part-of-speech tag.
- Tag:** The detailed part-of-speech tag.
- Dep:** Syntactic dependency, i.e. the relation between tokens.
- Shape:** The word `shape` – capitalization, punctuation, digits.
- is alpha:** Is the token an alpha character?
- is stop:** Is the token part of a stop list, i.e. the most common words of the language?

BREAK

Be back at 9 :35

Industrial-Strength Natural Language Processing

IN PYTHON

Get things done

spaCy is designed to help you do real work — to build real products, or gather real insights. The library respects your time, and tries to avoid wasting it. It's easy to install, and its API is simple and productive. We like to think of spaCy as the Ruby on Rails of Natural Language Processing.

GET STARTED

Blazing fast

spaCy excels at large-scale information extraction tasks. It's written from the ground up in carefully memory-managed Cython. Independent research in 2015 found spaCy to be the fastest in the world. If your application needs to process entire web dumps, spaCy is the library you want to be using.

FACTS & FIGURES

Deep learning

spaCy is the best way to prepare text for deep learning. It interoperates seamlessly with TensorFlow, PyTorch, scikit-learn, Gensim and the rest of Python's awesome AI ecosystem. With spaCy, you can easily construct linguistically sophisticated statistical models for a variety of NLP problems.

READ MORE

Lab

Project Group

Work on this sheet as a group: https://docs.google.com/document/d/1JpSFaH5x7_BOCrzOkhoUgzBAWs3JBeQeHDv_f-Fs1IY/edit?usp=sharing

PS2 discussion